



ADVANCED COMPUTER ARCHITECTURE



Khoa Khoa học và Kỹ thuật Máy tính
BM Kỹ thuật Máy tính

Trần Ngọc Thịnh

<http://www.cse.hcmut.edu.vn/~tnthinh>

©2013, dce

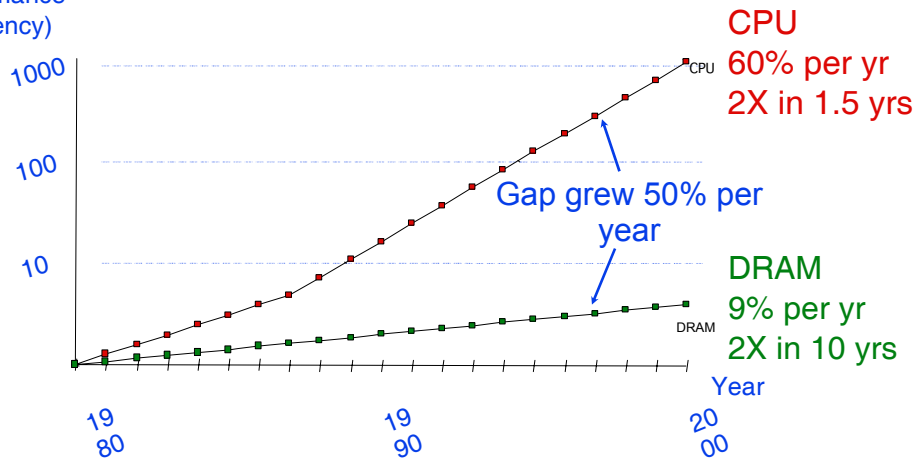
Memory Hierarchy Design



Since 1980, CPU has outpaced DRAM ...

Four-issue 2GHz superscalar accessing 100ns DRAM could execute 800 instructions during time for one memory access!

Performance
(1/latency)



3

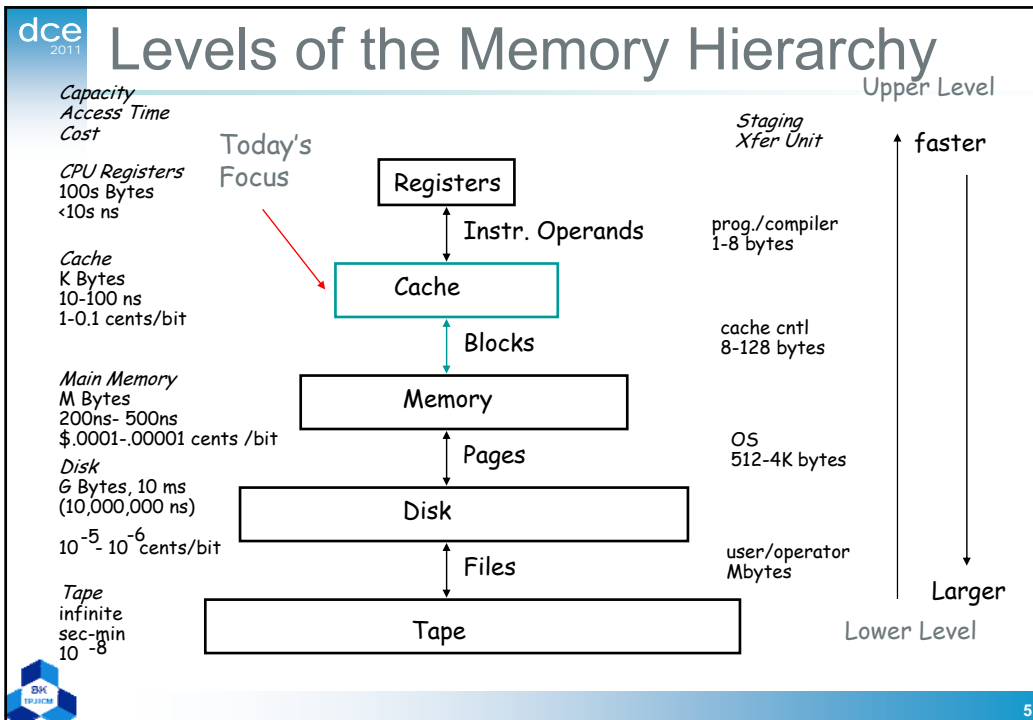
Processor-DRAM Performance Gap Impact

- To illustrate the performance impact, assume a single-issue pipelined CPU with CPI = 1 using non-ideal memory.
- Ignoring other factors, the minimum cost of a full memory access in terms of number of wasted CPU cycles:

Year	CPU speed MHZ	CPU cycle ns	Memory Access ns	<u>Minimum CPU memory stall cycles</u> or instructions wasted	
1986:	8	125	190	$190/125 - 1$	= 0.5
1989:	33	30	165	$165/30 - 1$	= 4.5
1992:	60	16.6	120	$120/16.6 - 1$	= 6.2
1996:	200	5	110	$110/5 - 1$	= 21
1998:	300	3.33	100	$100/3.33 - 1$	= 29
2000:	1000	1	90	$90/1 - 1$	= 89
2002:	2000	.5	80	$80/.5 - 1$	= 159
2004:	3000	.333	60	$60/.333 - 1$	= 179



4



dce 2011

Addressing the Processor-Memory Performance GAP

- **Goal:** Illusion of large, fast, cheap memory. Let programs address a memory space that scales to the disk size, at a speed that is usually as fast as register access
- **Solution:** Put smaller, faster “cache” memories between CPU and DRAM. Create a “memory hierarchy”.

6

Common Predictable Patterns

Two predictable properties of memory references:

- Temporal Locality: If a location is referenced, it is likely to be referenced again in the near future (e.g., loops, reuse).
- Spatial Locality: If a location is referenced it is likely that locations near it will be referenced in the near future (e.g., straightline code, array access).



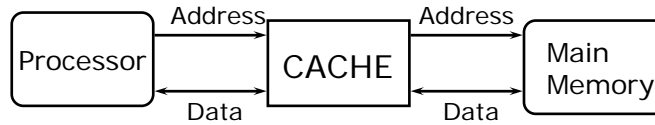
Caches

Caches exploit both types of predictability:

- Exploit **temporal locality** by remembering the contents of recently accessed locations.
- Exploit **spatial locality** by fetching blocks of data around recently accessed locations.



Simple view of cache



- The processor accesses the cache first
- **Cache hit:** Just use the data
- **Cache miss:** replace a block in cache by a block from main memory, use the data
- The data transferred between cache and main memory is in blocks, and controlled by independent hardware



Simple view of cache

- Hit rate: fraction of cache hit
- Miss rate: $1 - \text{Hit rate}$
- Miss penalty: Time to replace a block + time to deliver the data to the processor

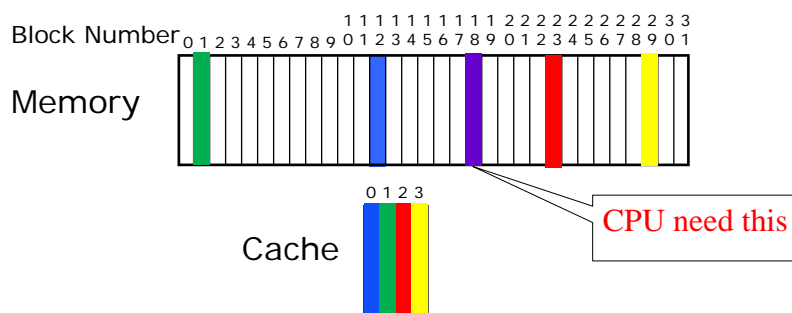


Simple view of cache

- Example: `For(i = 0; i < 10; i++) S = S + A[i];`
- No cache: At least 12 accesses to main memory (10 $A[i]$ and Read S , write S)
- With Cache: if $A[i]$ and S is in a single block (ex 32-bytes), 1 access to load block to cache, and 1 access to write block to main memory
- Access to S : Temporal Locality
- Access to $A[i]$: Spatial Locality ($A[i]$)



Replacement



- Cache cannot hold all blocks
- Replace a block by another that is currently needed by CPU

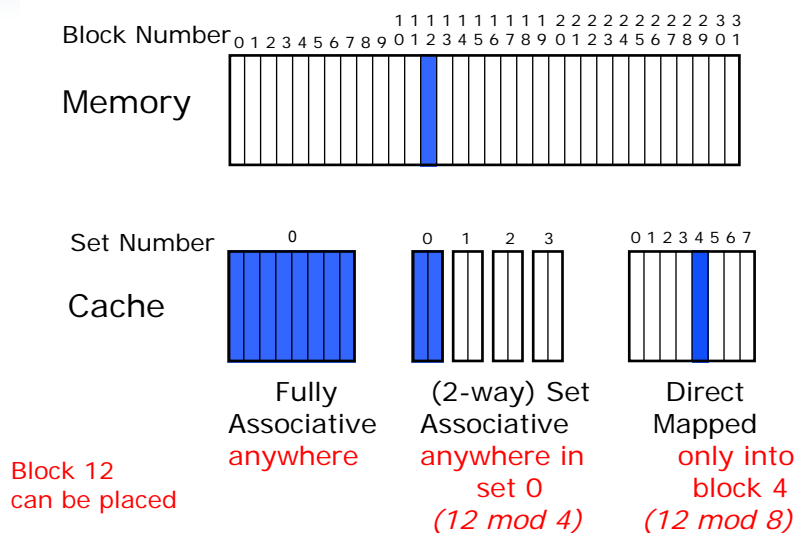


Basic Cache Design & Operation Issues

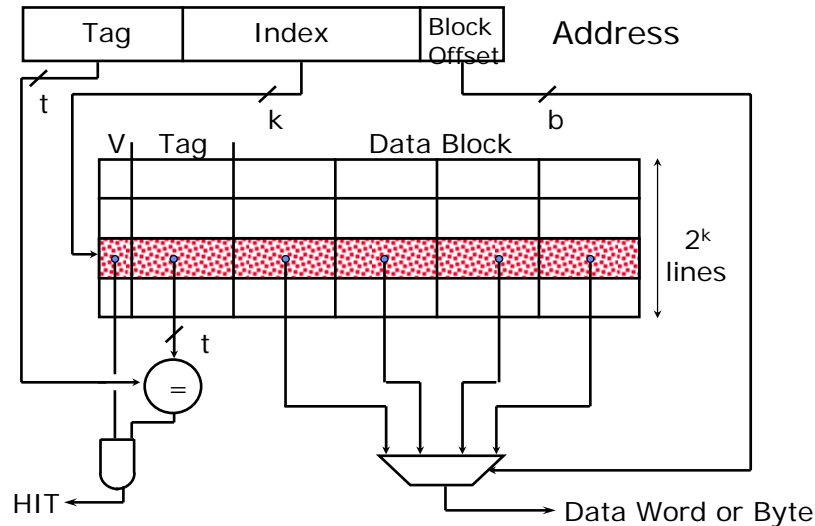
- Q1: Where can a block be placed cache?
(*Block placement strategy & Cache organization*)
 - Fully Associative, Set Associative, Direct Mapped.
- Q2: How is a block found if it is in cache?
(*Block identification*)
 - Tag/Block.
- Q3: Which block should be replaced on a miss?
(*Block replacement*)
 - Random, LRU, FIFO.
- Q4: What happens on a write?
(*Cache write policy*)
 - Write through, write back.



Q1: Where can a block be placed?



Direct-Mapped Cache

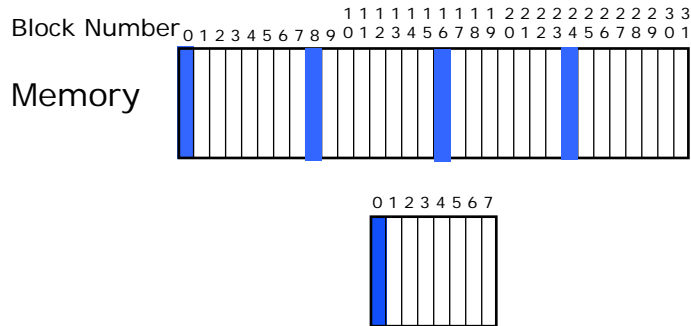


Direct-mapped Cache

- Address: N bits (2^N words)
- Cache has 2^k lines (blocks)
- Each line has 2^b words
- Block M is mapped to the line $M \% 2^k$
- Need $t = N - k - b$ **Tag** bits to identify mem. block
- Advantage: Simple
- Disadvantage: High miss rate
- What if CPU accesses block N_0 , N_1 and $N_0 \% 2^k = N_1 \% 2^k$?



Direct-mapped Cache



- Access N_0, N_1 where $N_0 \% 2^k = N_1 \% 2^k$
- Replace a block while there are many rooms available!



4KB Direct Mapped Cache Example

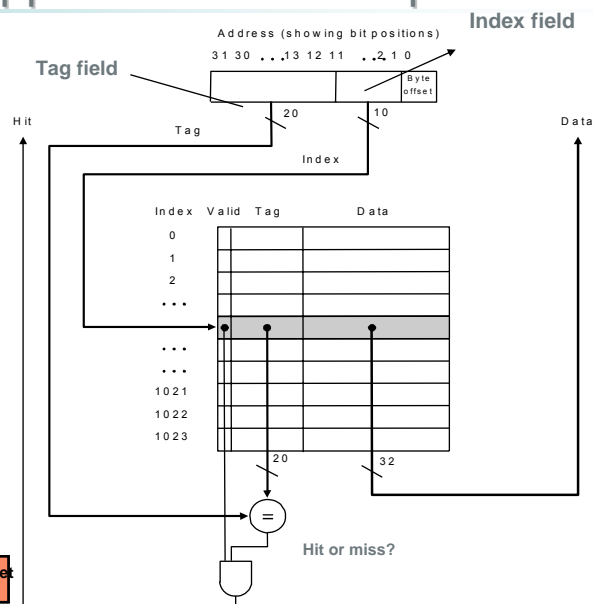
1K = 1024 Blocks
Each block = one **word**

Can cache up to
 2^{32} bytes = 4 GB
of memory

Mapping function:

Cache Block frame number =
(Block address) MOD (1024)

i.e. index field or
10 low bit of block address




Moi line 2^b word?

64KB Direct Mapped Cache Example

4K = 4096 blocks

Each block = four words = 16 bytes

Can cache up to 2^{32} bytes = 4 GB of memory




Larger cache blocks take better advantage of spatial locality and thus may result in a lower miss rate

Mapping Function: Cache Block frame number = (Block address) MOD (4096)
i.e. index field or 12 low bit of block address



Fully Associative Cache

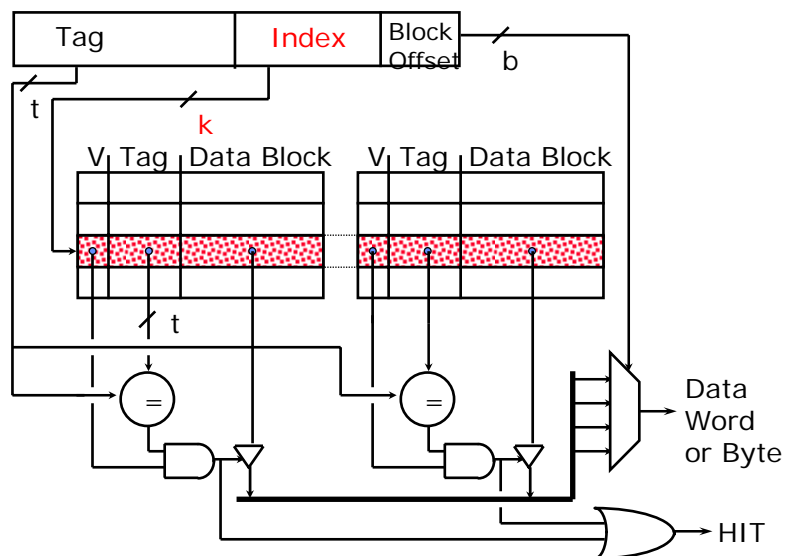


Fully associative cache

- CAM: Content Addressable Memory
- Each block can be mapped to any lines in cache
- Tag bit: $t = N - b$. Compared to Tag of all lines
- Advantage: replacement occurs only when no rooms available
- Disadvantage: resource consumption, delay by comparing many elements



Set-Associative Cache



W-way Set-associative Cache

- Balancing: Direct mapped cache vs Fully associative cache
- Cache has 2^k sets
- Each set has 2^w lines
- Block M is mapped to one of 2^w lines in set $M \% 2^k$
- Tag bit: $t = N - k - b$
- Currently: widely used (Intel, AMD, ...)



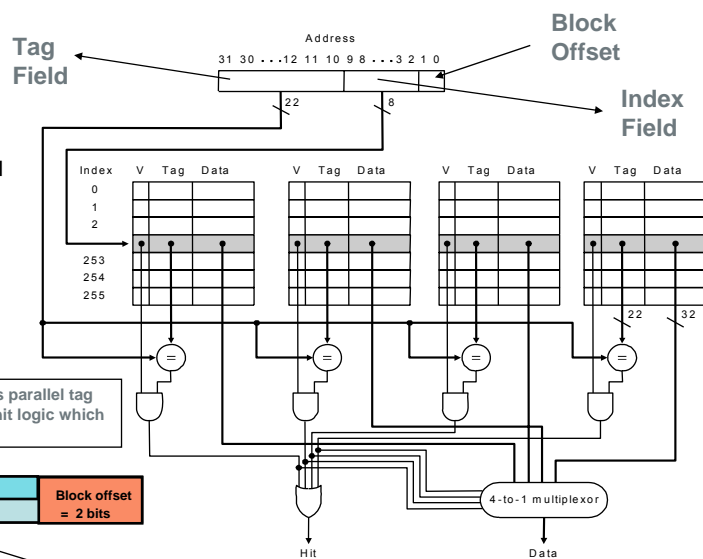
4K Four-Way Set Associative Cache: MIPS Implementation Example

1024 block frames
Each block = one word
4-way set associative
 $1024 / 4 = 256$ sets

Can cache up to
 2^{32} bytes = 4 GB
of memory

Set associative cache requires parallel tag matching and more complex hit logic which may increase hit time

Block Address = 30 bits	Block offset = 2 bits
Tag = 22 bits	Index = 8 bits



Mapping Function: Cache Set Number = $\text{index} = (\text{Block address}) \text{ MOD } (256)$



Q2: How is a block found?

- Index selects which set to look in
- Compare Tag to find block
- Increasing associativity shrinks index, expands tag. Fully Associative caches have no index field.
- Direct-mapped: 1-way set associative?
- Fully associative: 1 set?

Memory Address

Block Address		Block Offset
Tag	Index	



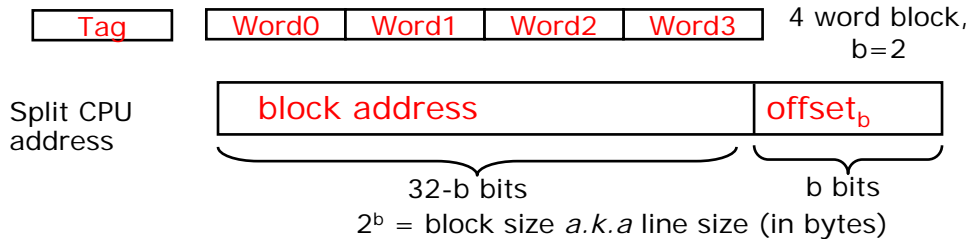
What causes a MISS?

- Three Major Categories of Cache Misses:
 - Compulsory Misses: first access to a block
 - Capacity Misses: cache cannot contain all blocks needed to execute the program
 - Conflict Misses: block replaced by another block and then later retrieved - (affects set assoc. or direct mapped caches)
 - Nightmare Scenario: ping pong effect!



Block Size and Spatial Locality

Block is unit of transfer between the cache and memory



Larger block size has distinct hardware advantages

- less tag overhead
- exploit fast burst transfers from DRAM
- exploit fast burst transfers over wide busses

What are the disadvantages of increasing block size?

Fewer blocks => more conflicts. Can waste bandwidth.



Q3: Which block should be replaced on a miss?

- Easy for Direct Mapped
- Set Associative or Fully Associative:
 - Random
 - Least Recently Used (LRU)
 - LRU cache state must be updated on every access
 - true implementation only feasible for small sets (2-way, 4-way)
 - pseudo-LRU binary tree often used for 4-8 way
 - First In, First Out (FIFO) a.k.a. Round-Robin
 - used in highly associative caches
- Replacement policy has a second order effect since replacement only happens on misses



Q4: What happens on a write?

- **Cache hit:**
 - *write through*: write both cache & memory
 - generally higher traffic but simplifies cache coherence
 - *write back*: write cache only
(memory is written only when the entry is evicted)
 - a dirty bit per block can further reduce the traffic
- **Cache miss:**
 - *no write allocate*: only write to main memory
 - *write allocate* (*aka fetch on write*): fetch into cache
- Common combinations:
 - write through and no write allocate (below example)
 - write back with write allocate (above Example)



Reading assignment 1

- Cache coherent problem in multicore systems
 - Identify the problem
 - Algorithms for multicore architectures
- Reference
 - eecs.wsu.edu/~cs460/cs550/cachecoherence.pdf
 - ...More on internet



Reading assignment 2

- Cache performance
 - Replacement policy (algorithms)
 - Optimization (Miss rate, penalty, ...)
- Reference
 - Hennessy - Patterson - Computer Architecture. A Quantitative
 - www2.lns.mit.edu/~avinatan/research/cache.pdf
 - ... More on internet